

# Big data in R with arrow :: CHEAT SHEET

APACHE  
ARROW

## About

Apache Arrow is a development platform for in-memory analytics. It contains a set of technologies that enable big data systems to process and move data fast.

The arrow R package integrates with dplyr and allows you to work with multiple storage formats as well as data in AWS S3 and other similar cloud storage systems.

## Installation

To install from CRAN,

```
install.packages("arrow")
```

To get all optional features enabled on Linux, you should set the environment variable `NOT_CRAN=true` prior to installing.

MacOS and Windows binary packages include all of these features by default.

Conda users can install with

```
conda install -c conda-forge --strict-channel-priority r-arrow
```

Visit [apache.arrow.org](https://arrow.apache.org) for the specifics to install the development version.

## Import

For **single files** you can do either:

```
read_parquet("gapminder.parquet")  
read_feather("gapminder.feather")
```

Arrow can also read large CSV and JSON files with excellent speed and efficiency:

```
read_csv_arrow("gapminder.csv")  
read_json_arrow("gapminder.json")
```

This reads data as `data.frame`.

To read **multiple Parquet/Feather files** from a directory you can specify a partitioning for filtering:

```
d <- open_dataset("nyc-taxi",  
  partitioning = c("year",  
    "month"))
```

This reads data as `ArrowObject` and needs posterior steps.

## Dplyr compatibility

Arrow in R shares most of the characteristics of SQL in R through `RPostgres` and other packages. The use of `collect()` converts Arrow-type objects into regular `data.frames` and allows you to use your data with your existing visualisation and analysis workflow.

If an operation is not yet implemented in Arrow, you can `collect()` and then do the operation on the `R data.frame`. This selection of data does need to fit into memory, but the whole dataset does not, you just need to use Arrow to select/filter down to something that does.

`dplyr` shall read exactly the required data fragments after you specified a partitioning without the need for an index as in SQL, here's an example of this:

```
d2 <- d %>%  
  filter(year == 2009,  
    month == 1) %>%  
  collect() %>%  
  group_by(year, month) %>%  
  summarise(mean_amount =  
    mean(total_amount))
```

This takes an `ArrowObject` and creates a `data.frame`.

## Export

To save without partitioning, you can use:

```
write_parquet(d2, "d2.parquet")  
write_feather(d2, "d2.feather")  
write_csv_arrow(d2, "d2.csv")
```

This is very similar to `write_csv()` from `readr`.